

# Machine Learning for Structured Prediction

Grzegorz Chrupała

National Centre for Language Technology  
School of Computing  
Dublin City University

NCLT Seminar



# Structured vs Non-structured prediction

In supervised learning we try to learn a function  $h : \mathcal{X} \rightarrow \mathcal{Y}$  where  $x \in \mathcal{X}$  are inputs and  $y \in \mathcal{Y}$  are outputs.

- Binary classification:  $\mathcal{Y} = \{-1, +1\}$
- Multiclass classification:  $\mathcal{Y} = \{1, \dots, K\}$  (finite set of labels)
- Regression:  $\mathcal{Y} = \mathbb{R}$
- In contrast, in structured prediction elements  $\mathcal{Y}$  are *complex*

The prediction is based on the *feature function*  $\Phi : \mathcal{X} \rightarrow \mathcal{F}$  where usually  $\mathcal{F} = \mathbb{R}^D$  ( $D$ -dimensional vector space)

# Structured prediction tasks

- Sequence labeling: e.g. POS tagging
- Parsing: given an input sequence, build a tree whose yield (leaves) are the elements in the sequence and whose structure obeys some grammar.
- Collective classification: E.g. relation learning problems, such as labeling web pages given link information.
- Bipartite matching: given a bipartite graph, find the best possible matching. e.g. word alignment in NLP and protein structure prediction in computational biology.

# Linear classifiers

- In binary classification, linear classifiers separate classes in a multidimensional input space with a hyperplane
- The classification function giving the output score is

$$y = f(\vec{x} \cdot \vec{w} + b) = f(\sum_{d \in D} w_d x_d + b)$$

where

- ▶  $\vec{x}$  is the input feature vector =  $\Phi(x)$ ,
- ▶  $\vec{w}$  is a real vector of feature weights,
- ▶  $b$  is the bias term, and
- ▶  $f$  maps the dot product of the two vectors to output (e.g.  $f(x) = -1$  if  $x < 0$ ,  $+1$  otherwise)
- ▶ The use of *kernels* permits to use linear classifiers for non-linearly separable problems such as the XOR function

# Linear classifier – hyperplanes




Figure: Three separating hyperplanes in 2-dimensional space (Wikipedia)

# Perceptron

- The Perceptron iterates  $I$  times through the  $1 \dots n$  training examples, updating the parameters (weights and bias) if the current example  $n$  is classified incorrectly:
  - ▶  $\vec{w} \leftarrow \vec{w} + y_n \vec{x}_n$
  - ▶  $b \leftarrow b + y_n$

where  $y_n$  is the output for the  $n^{th}$  training example

- Averaged Perceptron, which generalizes better to unseen examples, does *weight averaging*, i.e. the final weights returned are the average of all weight vectors used during the run of the algorithm

# Perceptron Algorithm

**Algorithm** AVERAGEDPERCEPTRON( $x_{1:N}, y_{1:N}, I$ )

```
1:  $\mathbf{w}_0 \leftarrow \langle 0, \dots, 0 \rangle, b_0 \leftarrow 0$ 
2:  $\mathbf{w}_a \leftarrow \langle 0, \dots, 0 \rangle, b_a \leftarrow 0$ 
3:  $c \leftarrow 1$ 
4: for  $i = 1 \dots I$  do
5:   for  $n = 1 \dots N$  do
6:     if  $y_n [\mathbf{w}_0^\top \Phi(x_n) + b_0] \leq 0$  then
7:        $\mathbf{w}_0 \leftarrow \mathbf{w}_0 + y_n \Phi(x_n), b_0 \leftarrow b_0 + y_n$ 
8:        $\mathbf{w}_a \leftarrow \mathbf{w}_a + c y_n \Phi(x_n), b_a \leftarrow b_a + c y_n$ 
9:     end if
10:     $c \leftarrow c + 1$ 
11:  end for
12: end for
13: return  $(\mathbf{w}_0 - \mathbf{w}_a/c, b_0 - b_a/c)$ 
```

Daume III (2006)

# Structured Perceptron

- For structured prediction the feature function is  $\Phi : \mathcal{X} \times \mathcal{Y} \rightarrow \mathcal{F}$  (second argument is the hypothesized output)
- Structured prediction algorithms need  $\Phi$  to admit tractable search
- The argmax problem:

$$\hat{y} = \underset{y \in \mathcal{Y}}{\operatorname{argmax}} \vec{w} \cdot \Phi(x, y)$$

- Structured Perceptron: for each example  $n$  wherever the predicted  $\hat{y}_n$  for  $x_n$  differs from  $y_n$  the weights are updated:
  - ▶  $\vec{w} \leftarrow \vec{w} + \Phi(x_n, y_n) - \Phi(x_n, \hat{y}_n)$
- Similar to plain Perceptron — key difference:  $\hat{y}$  is calculated using the arg max.
- Incremental Perceptron is a variant used in cases where arg max isn't available — a beam search algorithm is used instead



# Logistic regression / Maximum Entropy

- Conditional probability of class  $y$  is proportional to  $\exp f(x)$

$$\begin{aligned} p(y|x; \vec{w}, b) &= \frac{1}{Z_{x; \vec{w}, b}} \exp [y(\vec{w} \cdot \Phi(x) + b)] \\ &= \frac{1}{1 + \exp [-2y(\vec{w} \cdot \Phi(x) + b)]} \end{aligned}$$

- For training we try to find  $\vec{w}$  and  $b$  which maximize the likelihood of training data
- Gaussian prior used to penalize large weights
- Generalized Iterative Scaling algorithm used to solve the maximization problem

# Maximum Entropy Markov Models

- MEMMs are an extension of MaxEnt to sequence labeling
- Replace  $p(\text{observation}|\text{state})$  with  $p(\text{state}|\text{observation})$
- For first-order MEMM, the conditional distribution on the label  $y_n$  given full input  $x$ , the previous label  $y_{n-1}$ , feature function  $\Phi$  and weights  $\vec{w}$  is:

$$p(y_n|x, y_{n-1}; \vec{w}) = \frac{1}{Z_{x; y_{n-1}} \vec{w}} \exp [\vec{w} \cdot \Phi(x, y_n, y_{n-1})]$$

- When MEMM is trained *true* label  $y_{n-1}$  is used
- At prediction time, Viterbi is applied to solve  $\arg \max$
- *Predicted*  $y_{n-1}$  label is used in classifying  $n^{\text{th}}$  element of the sequence

# Conditional Random Fields

- Alternative generalization of MaxEnt to structured prediction
- Main difference:
  - ▶ while MEMM uses per state exponential models for the conditional probabilities of next states given the current state
  - ▶ CRF has a single exponential model for the joint probability of the whole sequence of labels
- The formulation is:

$$p(y|x; \vec{w}) = \frac{1}{Z_{x; \vec{w}}} \exp [\vec{w} \cdot \Phi(x, y)]$$
$$Z_{x; \vec{w}} = \sum_{y' \in \mathcal{Y}} \exp [\vec{w} \cdot \Phi(x, y')]$$

where elements  $y' \in \mathcal{Y}$  are incorrect outputs

# CRF - experimental comparison to HMM and MEMMs

<i>model</i>	<i>error</i>	<i>oov error</i>
HMM	5.69%	45.99%
MEMM	6.37%	54.61%
CRF	5.55%	48.05%
MEMM <sup>+</sup>	4.81%	26.99%
CRF <sup>+</sup>	4.27%	23.76%

<sup>+</sup>Using spelling features

*Figure 4.* Per-word error rates for POS tagging on the Penn tree-bank, using first-order models trained on 50% of the 1.1 million word corpus. The oov rate is 5.45%.

J. Lafferty, A. McCallum, F. Pereira - Proc. 18th International Conf. on Machine Learning, 2001

# Large Margin Models

- Choose parameter settings which maximize the distance between the hyperplane separating classes and the nearest data points on either side.
- Support Vector Machines. For data separable with margin 1:

$$\text{minimize}_{\vec{w}, b} \quad \frac{1}{2} \|\vec{w}\|^2$$

subject to

$$\forall n, \quad y_n [\vec{w} \cdot \Phi(x_n) + b] \geq 1$$

- Soft-margin formulation: don't enforce perfect classification of training examples
- The *slack* variables  $\xi$  measures how far an example is from the correct hard margin constraint


$$\underset{\vec{w}, b}{\text{minimize}} \quad \frac{1}{2} \|\vec{w}\|^2 + C \sum_{n=1}^N \xi_n$$

subject to

$$\begin{aligned} \forall n, \quad y_n [\vec{w} \cdot \Phi(x_n) + b] &\geq 1 - \xi_n \\ \xi_n &\geq 0 \end{aligned}$$

- The parameter  $C$  trades off fitting the training data and a small weight vector

# Margins and Support Vectors



**Figure:** SVM maximum-margin hyperplanes for binary classification. Examples along the hyperplanes are support vectors (from Wikipedia)

# Maximum Margin models for Structured Prediction

- Maximum Margin Markov Networks ( $M^3N$ )
- The difference in score (under loss function  $l$ ) between true output  $y$  and any incorrect output  $\hat{y}$  is at least  $l(x, y, \hat{y})$  —  $M^3N$  scales the margin to be proportional to loss.
- Support Vector Machines for Interdependent and Structured Outputs:  $SVM^{struct}$
- Instead of scaling the margin, scale the slack variables by the loss
- The two approaches differ in the optimization techniques used ( $SVM^{struct}$  constrains the loss function less but is more difficult to optimize).



# Summary of structured prediction algorithms

	Loss			Features			Efficient	Easy to Implement
	0/1	Hamming	Any	argmax and sum	argmax only	Neither		
Structured Perceptron	✓				✓		✓	✓
Conditional Random Field	✓			✓				—
Max-margin Markov Network	✓	✓			✓			
SVM for Structured Outputs	✓	✓			✓			
Reranking	✓	✓	✓			✓	—	—

Figure from Daume III (2006)

# Other approaches to structured outputs

- Searn (Search + Learn)
- Analogical learning
- Miscellaneous hacks:
  - ▶ Reranking
  - ▶ Stacking
  - ▶ Class n-grams

# References

- M. Collins, 2002, Discriminative Training Methods for Hidden Markov Models: Theory and Experiments with Perceptron Algorithms, EMNLP-02
- McCallum, Andrew, Dayne Freitag, and Fernando Pereira. 2000. Maximum entropy Markov models for information extraction and segmentation. ICML
- J. Lafferty, A. McCallum, and F. Pereira. 2001. Conditional random elds: Probabilistic models for segmenting and labeling sequence data. ICML
- B. Taskar, C. Guestrin, and D. Koller. 2003. Max-margin Markov networks. In Advances in Neural Information Processing Systems (NIPS).
- I. Tsochantaridis, T. Hofmann, T. Joachims, and Y. Altun. 2005. Large margin methods for structured and interdependent output variables. JMLR
- Hal Daumé III, 2006, *Practical Structured Learning Techniques for Natural Language Processing*, PhD thesis